

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler

```

```

import warnings
warnings.filterwarnings('ignore')

```

```

data=pd.read_csv('Boston housing dataset.csv')

```

```

data.head

```

<bound method NDFrame.head of						CRIM	ZN	INDUS	CHAS	NOX	
RM	AGE	DIS	RAD	TAX	\						
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
..	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	NaN	36.2
..	...	...	...	...
501	21.0	391.99	NaN	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9

```
504      21.0  393.45   6.48  22.0
505      21.0  396.90   7.88  11.9
```

```
[506 rows x 14 columns]>
```

```
data.shape
```

```
(506, 14)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 506 entries, 0 to 505
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	CRIM	486 non-null	float64
1	ZN	486 non-null	float64
2	INDUS	486 non-null	float64
3	CHAS	486 non-null	float64
4	NOX	506 non-null	float64
5	RM	506 non-null	float64
6	AGE	486 non-null	float64
7	DIS	506 non-null	float64
8	RAD	506 non-null	int64
9	TAX	506 non-null	int64
10	PTRATIO	506 non-null	float64
11	B	506 non-null	float64
12	LSTAT	486 non-null	float64
13	MEDV	506 non-null	float64

```
dtypes: float64(12), int64(2)
```

```
memory usage: 55.5 KB
```

```
data.nunique()
```

CRIM	484
ZN	26
INDUS	76
CHAS	2
NOX	81
RM	446
AGE	348
DIS	412
RAD	9
TAX	66
PTRATIO	46
B	357
LSTAT	438
MEDV	229

```
dtype: int64
```

```

data.CHAS.unique()
array([ 0., nan, 1.])

data.ZN.unique()
array([ 18. ,  0. , 12.5, 75. , 21. , 90. , 85. , 100. , 25. ,
        17.5, 80. , nan, 28. , 45. , 60. , 95. , 82.5, 30. ,
        22. , 20. , 40. , 55. , 52.5, 70. , 34. , 33. , 35. ])

data.isnull().sum()
CRIM      20
ZN         20
INDUS     20
CHAS      20
NOX        0
RM         0
AGE       20
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT     20
MEDV       0
dtype: int64

data.duplicated().sum()
0

np.int64(0)
0

df=data.copy()

df['CRIM'].fillna(df['CRIM'].mean(),inplace=True)
df['ZN'].fillna(df['ZN'].mean(),inplace=True)
df['CHAS'].fillna(df['CHAS'].mode()[0],inplace=True)
df['INDUS'].fillna(df['INDUS'].mean(),inplace=True)
df['AGE'].fillna(df['AGE'].median(),inplace=True)
df['LSTAT'].fillna(df['LSTAT'].median(),inplace=True)

df.isnull().sum()
CRIM      0
ZN         0
INDUS     0
CHAS      0
NOX        0

```

```
RM          0
AGE         0
DIS         0
RAD         0
TAX         0
PTRATIO     0
B           0
LSTAT       0
MEDV        0
dtype: int64
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	11.43	36.2

```
df['CHAS']=df['CHAS'].astype('int')
```

```
df.describe().T
```

	count	mean	std	min	25%
CRIM	506.0	3.611874	8.545770	0.00632	0.083235
ZN	506.0	11.211934	22.921051	0.00000	0.000000
INDUS	506.0	11.083992	6.699165	0.46000	5.190000
CHAS	506.0	0.067194	0.250605	0.00000	0.000000
NOX	506.0	0.554695	0.115878	0.38500	0.449000
RM	506.0	6.284634	0.702617	3.56100	5.885500

6.20850					
AGE	506.0	68.845850	27.486962	2.90000	45.925000
76.80000					
DIS	506.0	3.795043	2.105710	1.12960	2.100175
3.20745					
RAD	506.0	9.549407	8.707259	1.00000	4.000000
5.00000					
TAX	506.0	408.237154	168.537116	187.00000	279.000000
330.00000					
PTRATIO	506.0	18.455534	2.164946	12.60000	17.400000
19.05000					
B	506.0	356.674032	91.294864	0.32000	375.377500
391.44000					
LSTAT	506.0	12.664625	7.017219	1.73000	7.230000
11.43000					
MEDV	506.0	22.532806	9.197104	5.00000	17.025000
21.20000					

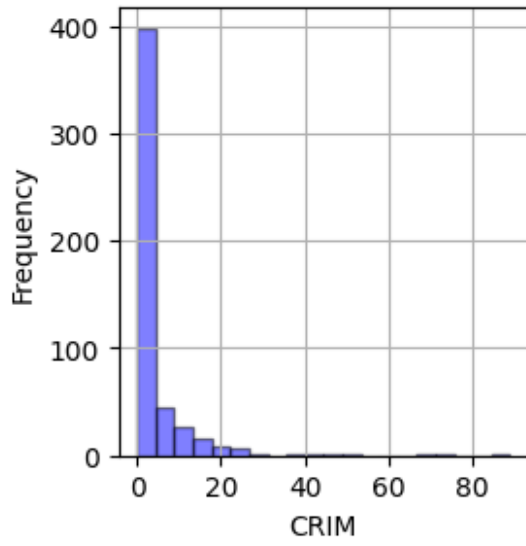
	75%	max
CRIM	3.611874	88.9762
ZN	11.211934	100.0000
INDUS	18.100000	27.7400
CHAS	0.000000	1.0000
NOX	0.624000	0.8710
RM	6.623500	8.7800
AGE	93.575000	100.0000
DIS	5.188425	12.1265
RAD	24.000000	24.0000
TAX	666.000000	711.0000
PTRATIO	20.200000	22.0000
B	396.225000	396.9000
LSTAT	16.570000	37.9700
MEDV	25.000000	50.0000

```


for i in df.columns:
    plt.figure(figsize=(6,3))
    plt.subplot(1,2,1)
    df[i].hist(bins=20,alpha=0.5,color='b',edgecolor='black')
    plt.title(f'Histogram of {i}')
    plt.xlabel(i)
    plt.ylabel('Frequency')
    plt.subplot(1,2,2)
    plt.boxplot(df[i],vert=False)
    plt.title(f'Boxplot of {i}')
    plt.show()

```


Histogram of CRIM




Boxplot of CRIM




Histogram of ZN




Boxplot of ZN




### Histogram of INDUS




### Boxplot of INDUS




### Histogram of CHAS




### Boxplot of CHAS




### Histogram of NOX




### Boxplot of NOX



### Histogram of RM




### Boxplot of RM







### Histogram of AGE




### Boxplot of AGE



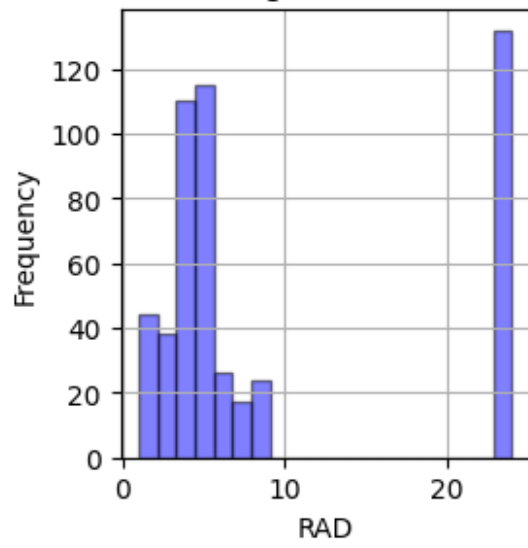
### Histogram of DIS




### Boxplot of DIS



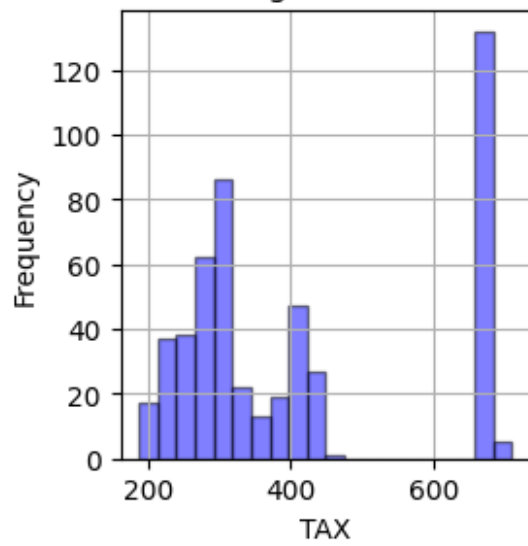
### Histogram of RAD




### Boxplot of RAD




### Histogram of TAX




### Boxplot of TAX




### Histogram of PTRATIO




### Boxplot of PTRATIO







### Histogram of B




### Boxplot of B





```
corr=df.corr(method='pearson')
plt.figure(figsize=(10,8))
sns.heatmap(corr,annot=True,cmap="coolwarm",fmt=".2f",linewidths=0.5)
plt.xticks(rotation=90,ha='right')
plt.yticks(rotation=0)
plt.title("Correlation Matrix Heatmap")
plt.show()
```



```

X=df.drop('MEDV',axis=1)
y=df['MEDV']

scale=StandardScaler()
X_scaled=scale.fit_transform(X)

X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=0.2,random_state=42)

model=LinearRegression()
model.fit(X_train,y_train)

LinearRegression()

y_pred=model.predict(X_test)
y_pred

```

```
array([28.99719439, 36.56606809, 14.51022803, 25.02572187,
18.42885474,
      23.02785726, 17.95437605, 14.5769479 , 22.14430832,
20.84584632,
      25.15283588, 18.55925182, -5.69168071, 21.71242445,
19.06845707,
      25.94275348, 19.70991322,  5.85916505, 40.9608103 ,
17.21528576,
      25.36124981, 30.26007975, 11.78589412, 23.48106943,
17.35338161,
      15.13896898, 21.61919056, 14.51459386, 23.17246824,
19.40914754,
      22.56164985, 25.21208496, 25.88782605, 16.68297496,
16.44747174,
      16.65894826, 31.10314158, 20.25199803, 24.38567686,
23.09800032,
      14.47721796, 32.36053979, 43.01157914, 17.61473728,
27.60723089,
      16.43366912, 14.25719607, 26.0854729 , 19.75853278,
30.15142187,
      21.01932313, 33.72128781, 16.39180467, 26.36438908,
39.75793372,
      22.02419633, 18.39453126, 32.81854401, 25.370573 ,
12.82224665,
      22.76128341, 30.73955199, 31.34386371, 16.27681305,
20.36945226,
      17.23156773, 20.15406451, 26.15613066, 30.92791361,
11.42177654,
      20.89590447, 26.58633798, 11.01176073, 12.76831709,
23.73870867,
      6.37180464, 21.6922679 , 41.74800223, 18.64423785,
8.82325704,
      20.96406016, 13.20179007, 20.99146149,  9.17404063,
23.0011185 ,
      32.41062673, 18.99778065, 25.56204885, 28.67383635,
19.76918944,
      25.94842754,  5.77674362, 19.514431 , 15.22571165,
10.87671123,
      20.08359505, 23.77725749,  0.05985008, 13.56333825,
16.1215622 ,
      22.74200442, 24.36218289])
```

```
mse=mean_squared_error(y_test,y_pred)
```

```
rmse=np.sqrt(mse)
```

```
r2=r2_score(y_test,y_pred)
```

```
print(f'Mean Squared Error:{mse}')  
print(f'Root Mean Squared Error:{rmse}')
```

```
print(f'R-Squared:{r2}')
```

Mean Squared Error:24.944071172175573

Root Mean Squared Error:4.99440398567993

R-Squared:0.6598556613717497